



Four Things to Know about Reliable Spark Streaming

Dean Wampler, Typesafe
Tathagata Das, Databricks

Agenda for today

- The Stream Processing Landscape
- How Spark Streaming Works - A Quick Overview
- Features in Spark Streaming that Help Prevent Data Loss
- Design Tips for Successful Streaming Applications

The Stream Processing Landscape

Stream Processors



STORM

samza

Spark

akka

Stream Storage



kafka



amazon
web services™

S3



hadoop
HDFS



riak



databricks™

Typesafe

Stream Sources

~~Ø~~MQ

MQTT



How Spark Streaming Works: A Quick Overview

Spark Streaming

Scalable, fault-tolerant stream processing system

High-level API

joins, windows, ...
often 5x less code

Fault-tolerant

Exactly-once semantics,
even for stateful ops

Integration


Integrates with MLlib, SQL,
DataFrames, GraphX



Spark Streaming

Receivers receive data streams and chop them up into batches

Spark processes the batches and pushes out the results



Word Count with Kafka

```
val context = new StreamingContext(conf, Seconds(1))
```

entry point of streaming
functionality

```
val lines = KafkaUtils.createStream(context, ...)
```

create DStream
from Kafka data

Word Count with Kafka

```
val context = new StreamingContext(conf, Seconds(1))
```

```
val lines = KafkaUtils.createStream(context, ...)
```

```
val words = lines.flatMap(_.split(" "))
```

split lines into words

Word Count with Kafka

```
val context = new StreamingContext(conf, Seconds(1))
```

```
val lines = KafkaUtils.createStream(context, ...)
```

```
val words = lines.flatMap(_.split(" "))
```

```
val wordCounts = words.map(x => (x, 1))
```

```
.reduceByKey(_ + _)
```

count the words

```
wordCounts.print()
```

print some counts on screen

```
context.start()
```

start receiving and
transforming the data

Word Count with Kafka

```
object WordCount {  
  def main(args: Array[String]) {  
    val context = new StreamingContext(new SparkConf(), Seconds(1))  
    val lines = KafkaUtils.createStream(context, ...)  
    val words = lines.flatMap(_.split(" "))  
    val wordCounts = words.map(x => (x,1)).reduceByKey(_ + _)  
    wordCounts.print()  
    context.start()  
    context.awaitTermination()  
  }  
}
```

Features in Spark Streaming that Help Prevent Data Loss

A Deeper View of Spark Streaming

Any Spark Application

User code runs in
the driver process

Spark
Driver



YARN / Mesos /
Spark Standalone
cluster

Any Spark Application

User code runs in
the driver process



Spark
Executor



Driver launches
executors in
cluster

Spark
Executor

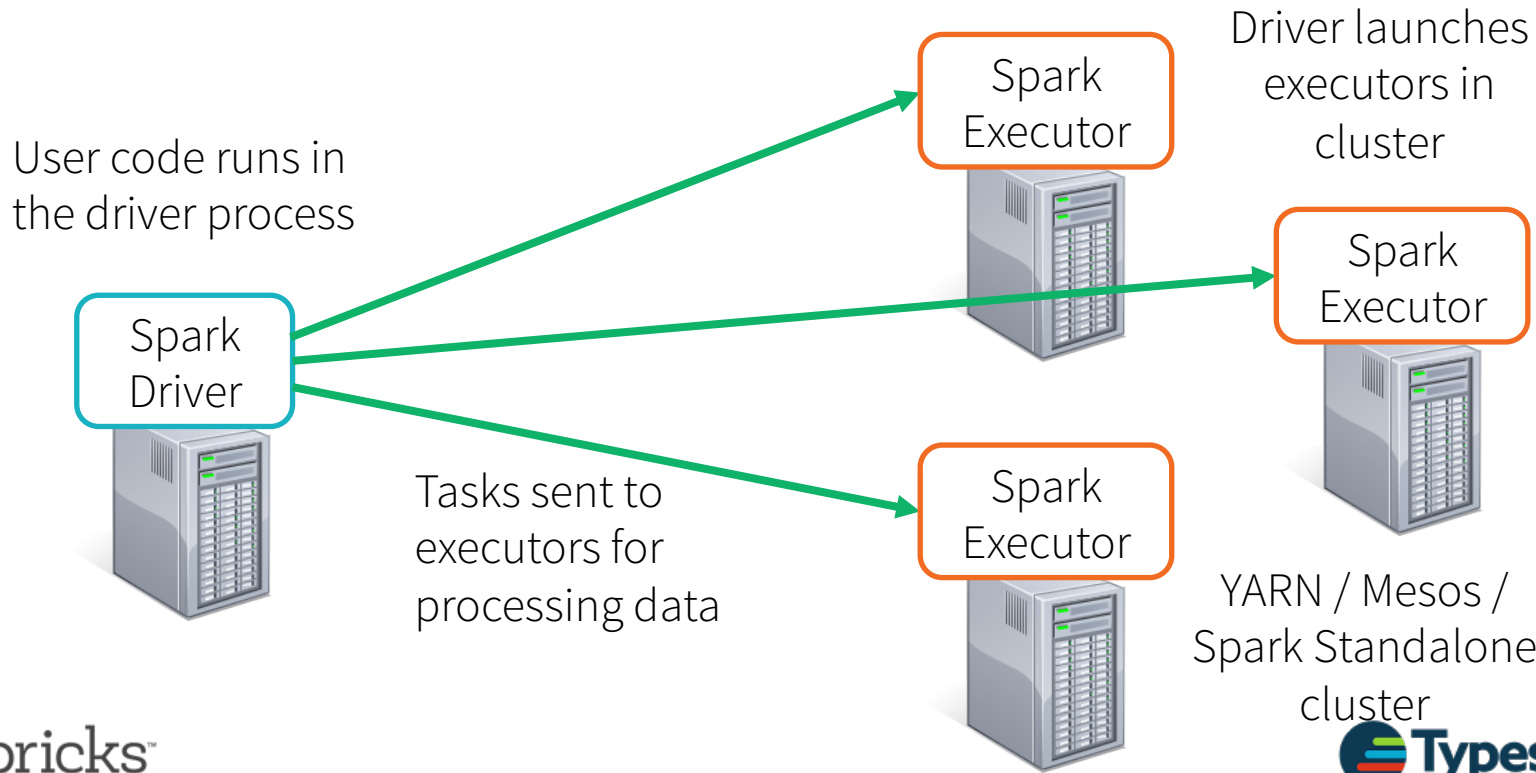


Spark
Executor

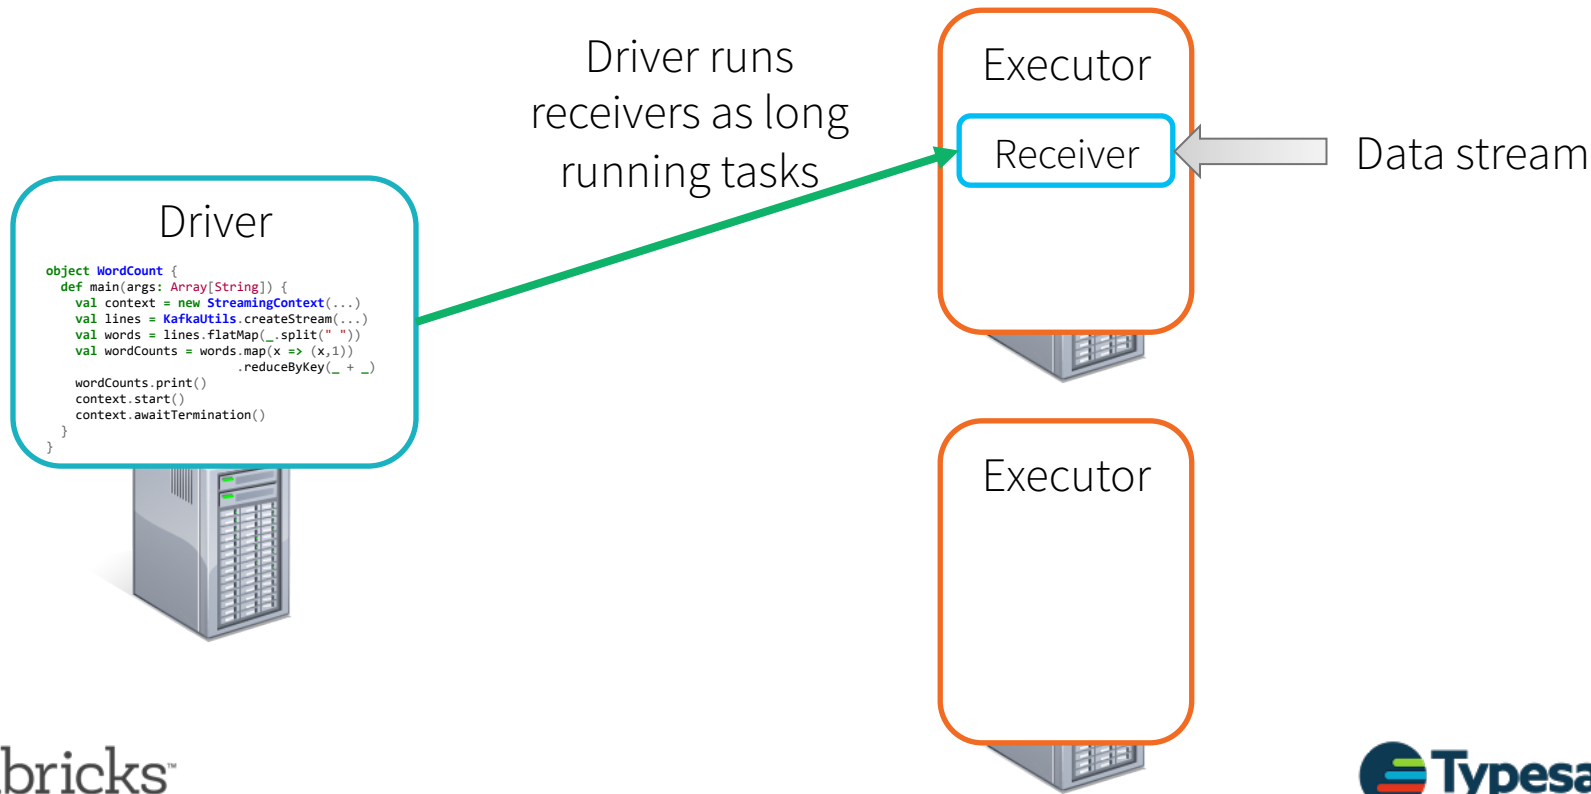


YARN / Mesos /
Spark Standalone
cluster

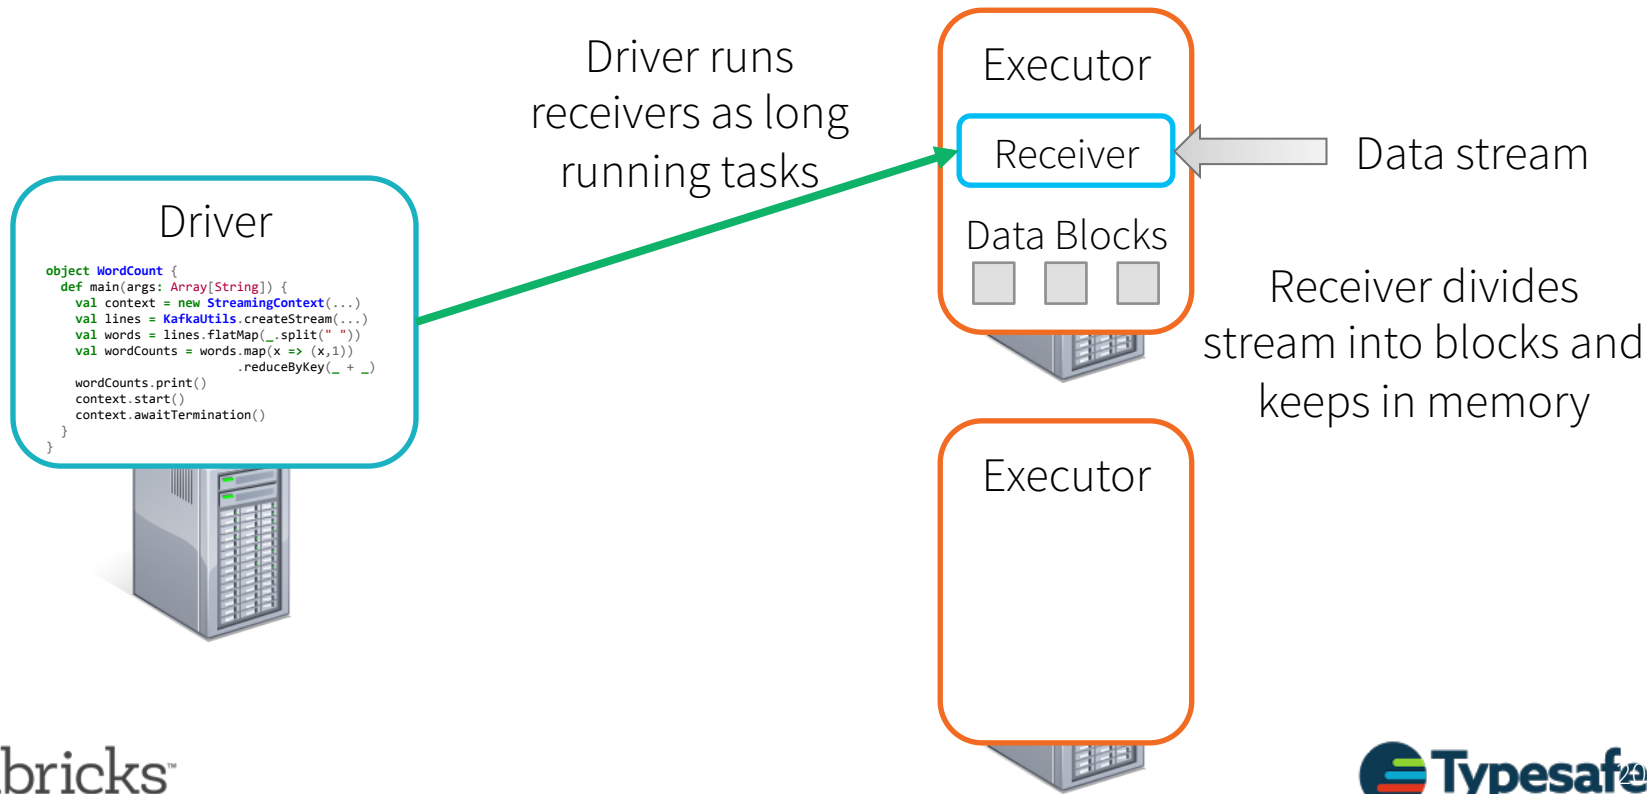
Any Spark Application



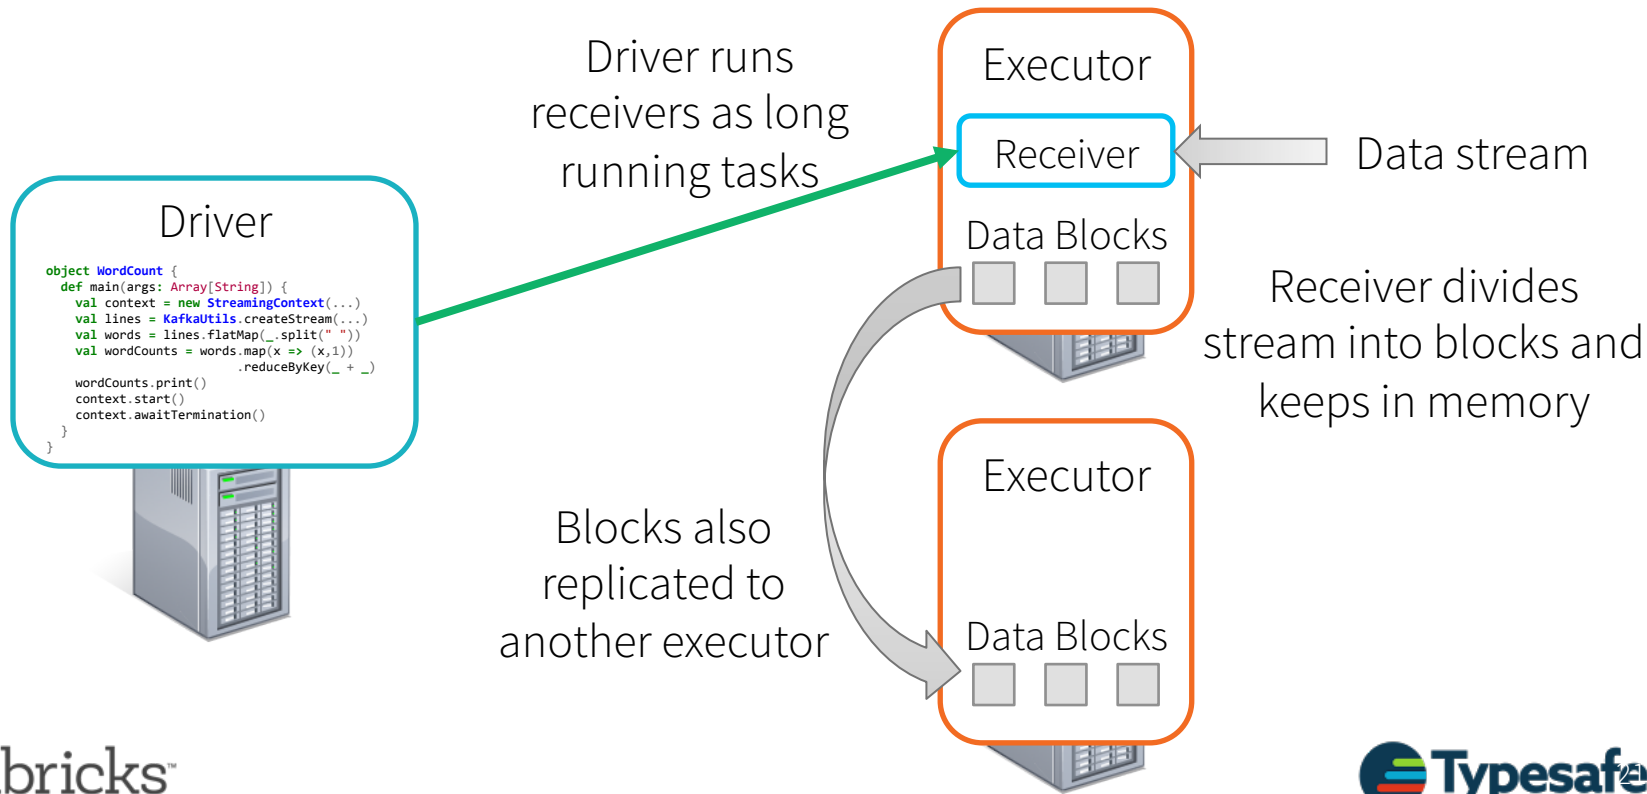
Spark Streaming Application: Receive data



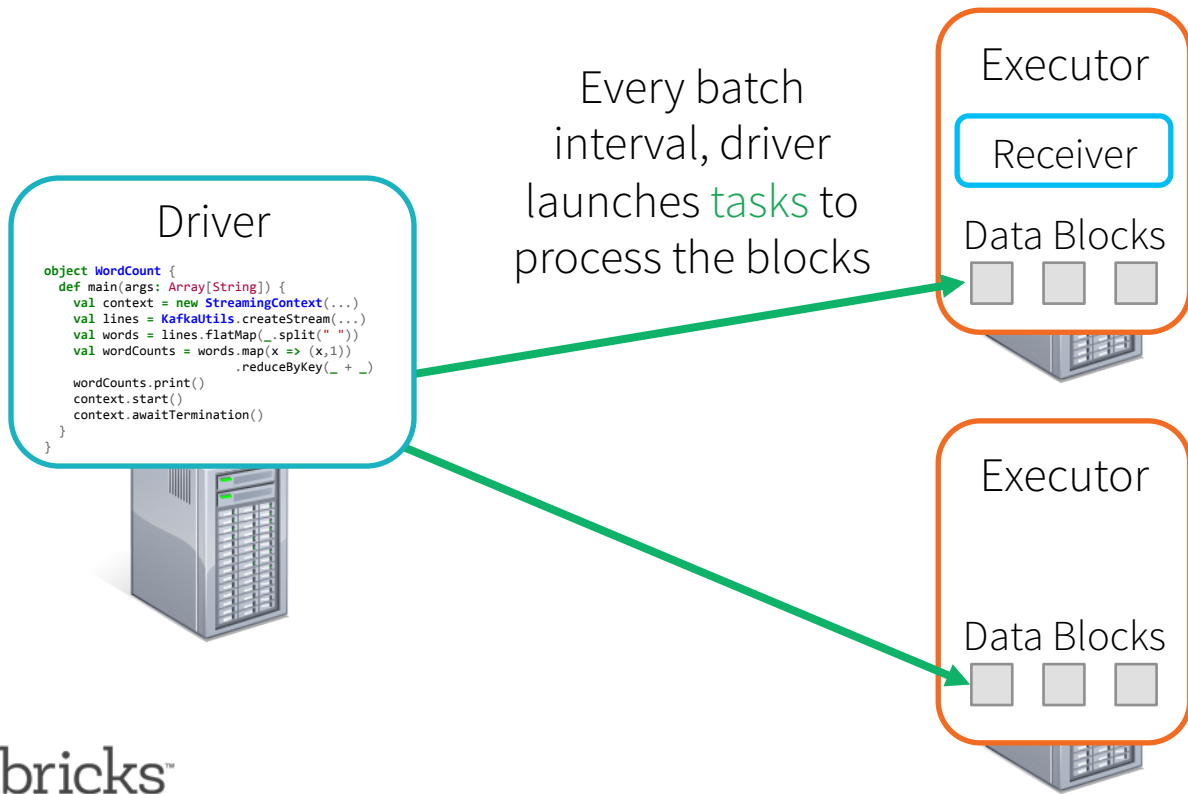
Spark Streaming Application: Receive data



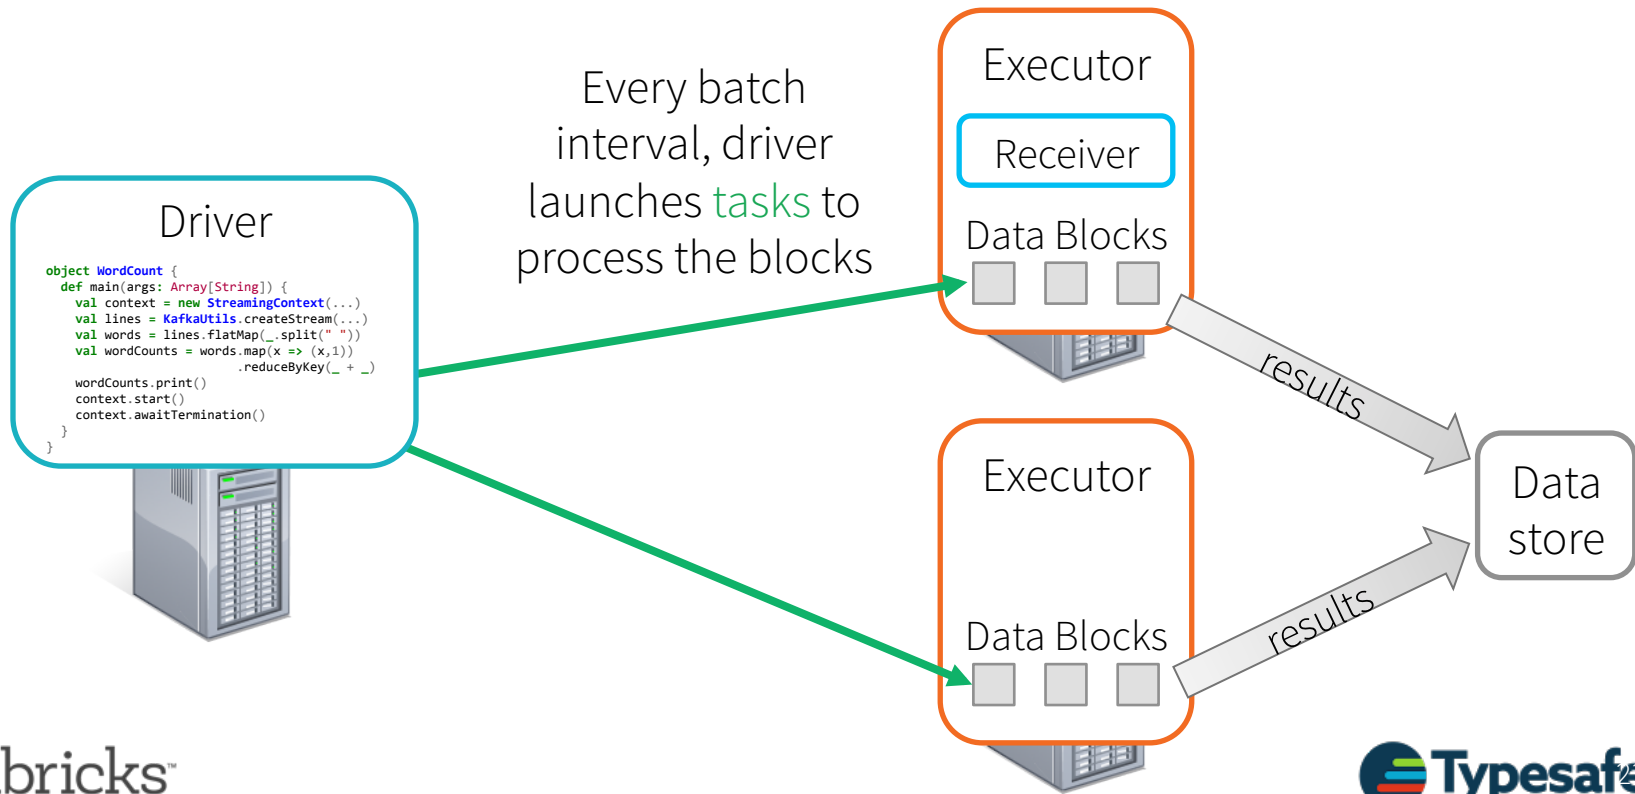
Spark Streaming Application: Receive data



Spark Streaming Application: Process data



Spark Streaming Application: Process data



Fault Tolerance and Reliability

Failures? Why care?

Many streaming applications need zero data loss guarantees despite any kind of failures in the system

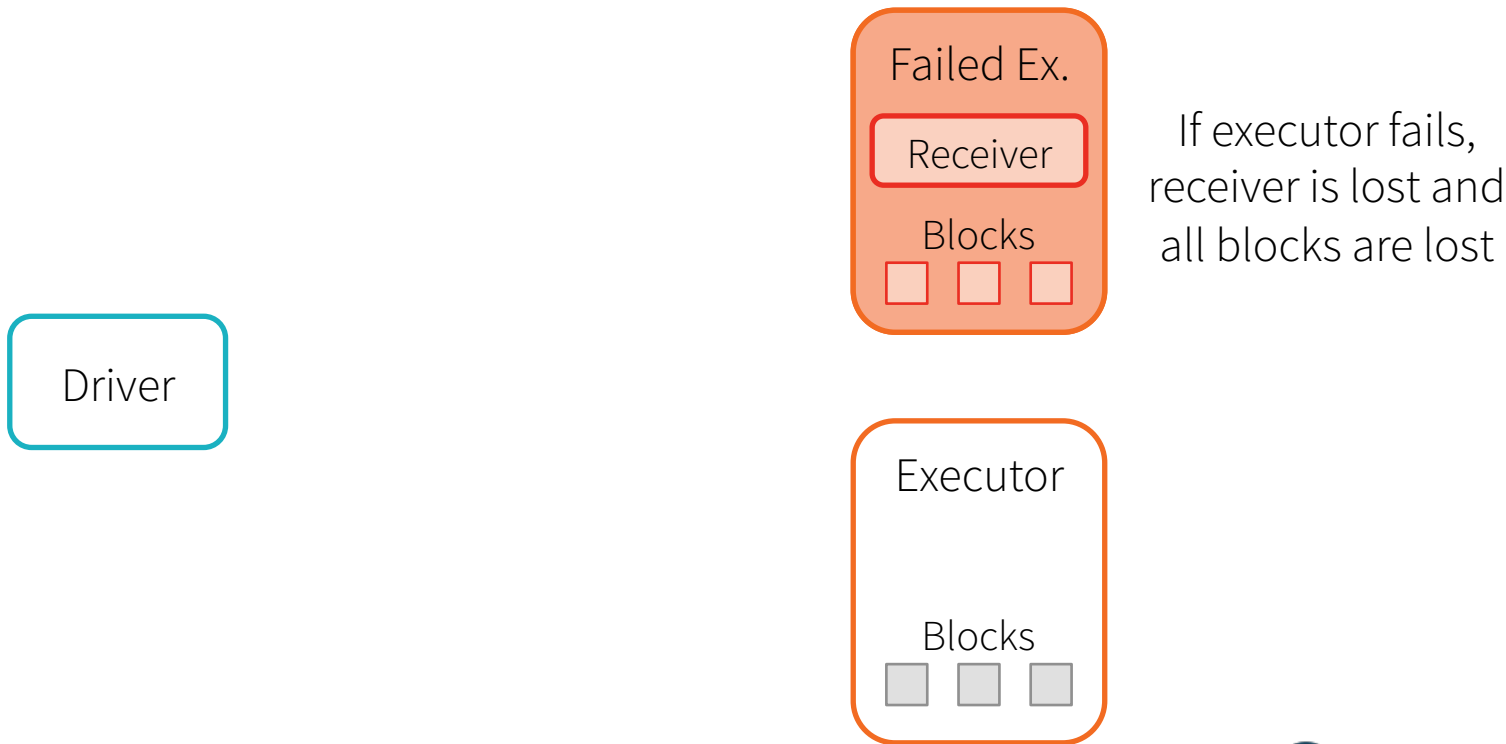
At least once guarantee – every record processed at least once

Exactly once guarantee – every record processed exactly once

Different kinds of failures – **executor** and **driver**

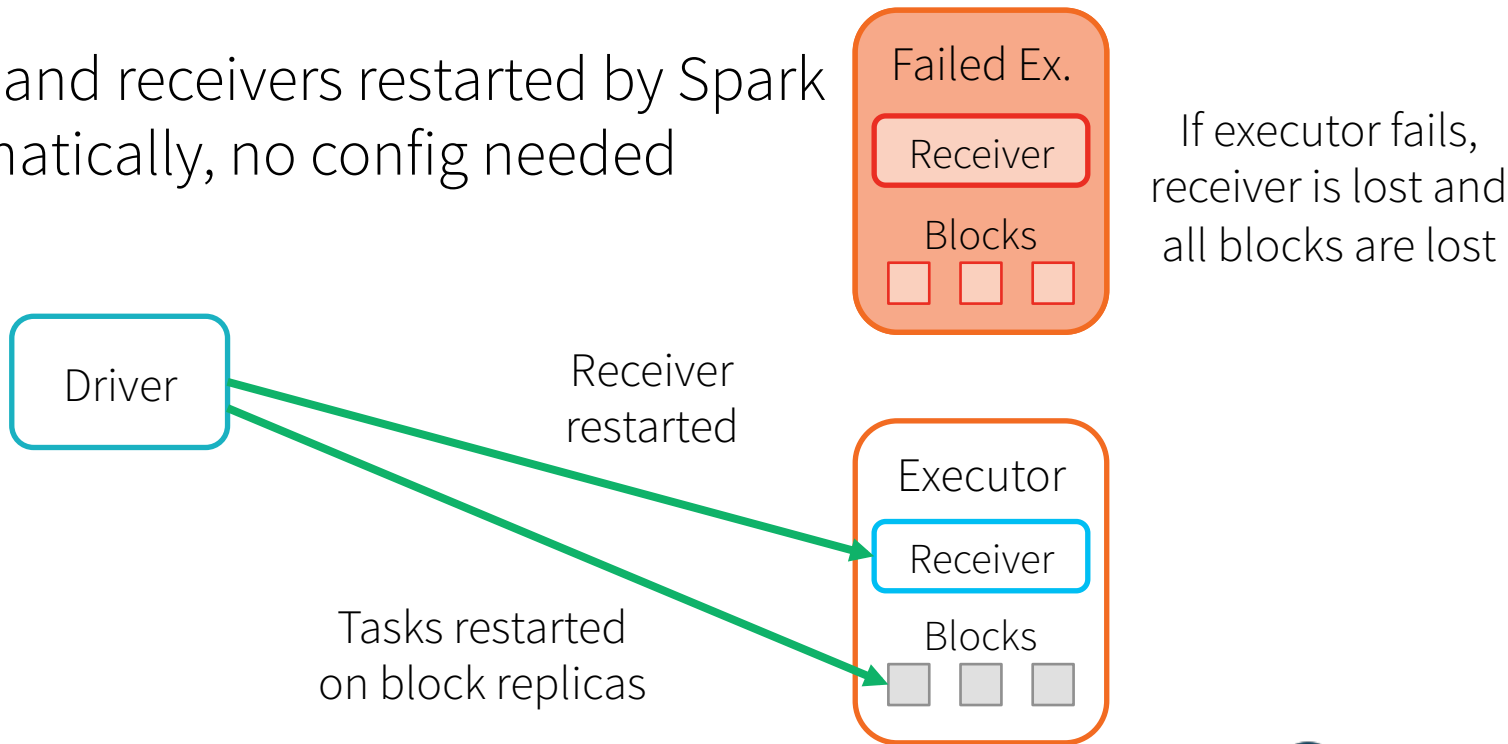
Some failures and guarantee requirements need additional configurations and setups

What if an executor fails?



What if an executor fails?

Tasks and receivers restarted by Spark automatically, no config needed



What if the driver fails?

Failed
Driver

When the driver
fails, all the
executors fail

All computation,
all received
blocks are lost

How do we
recover?

Failed Ex.

Receiver

Blocks



Failed
Executor




Blocks



Recovering Driver w/ DStream Checkpointing

DStream Checkpointing:

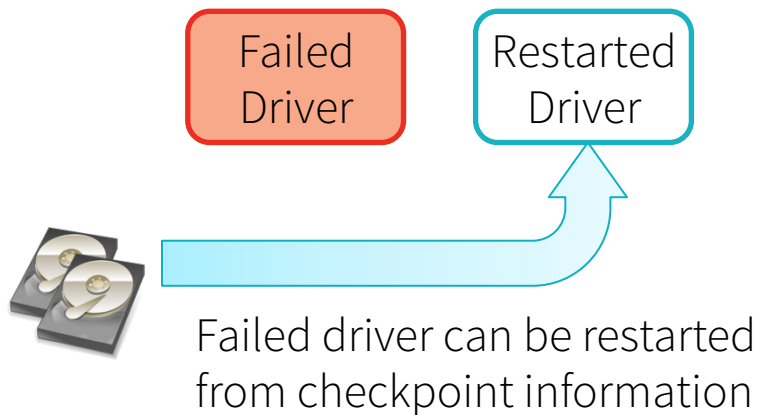
Periodically save the DAG of DStreams to fault-tolerant storage



Recovering Driver w/ DStream Checkpointing

DStream Checkpointing:

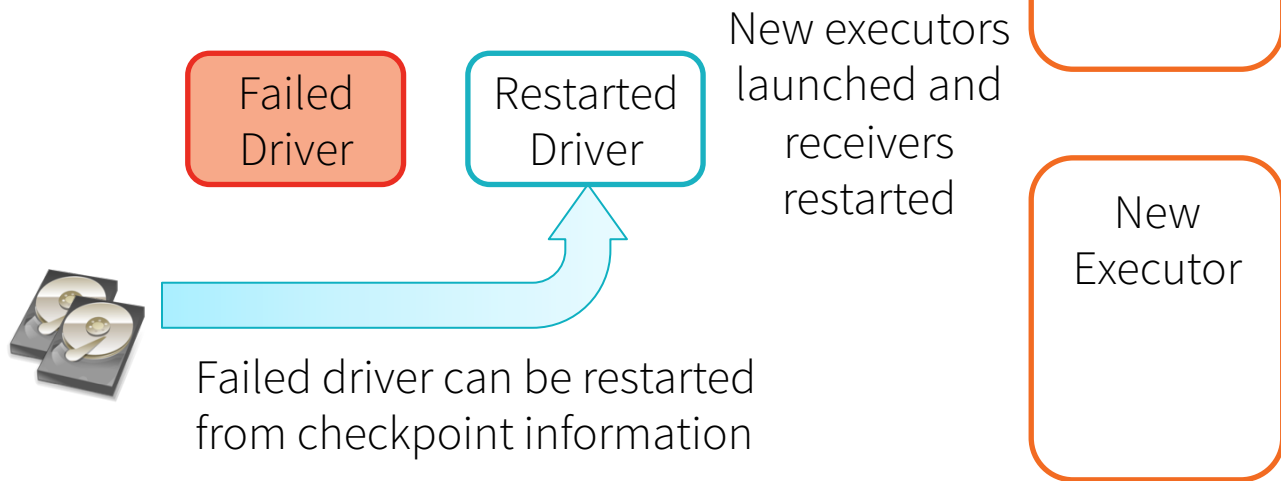
Periodically save the DAG of DStreams to fault-tolerant storage



Recovering Driver w/ DStream Checkpointing

DStream Checkpointing:

Periodically save the DAG of DStreams to fault-tolerant storage



Recovering Driver w/ DStream Checkpointing

1. Configure automatic driver restart
All cluster managers support this
2. Set a checkpoint directory in a HDFS-compatible file system
`streamingContext.checkpoint(hdfsDirectory)`
3. Slightly restructure of the code to use checkpoints for recovery

Configuring Automatic Driver Restart

Spark Standalone – Use spark-submit with “cluster” mode and “--supervise”

See <http://spark.apache.org/docs/latest/spark-standalone.html>

YARN – Use spark-submit in “cluster” mode

See YARN config “yarn.resourcemanager.am.max-attempts”

Mesos – Marathon can restart applications or use the “--supervise” flag.

Restructuring code for Checkpointing

Create
+
Setup

```
val context = new StreamingContext(...)
val lines = KafkaUtils.createStream(...)
val words = lines.flatMap(...)
...
```

Start

```
context.start()
```

Restructuring code for Checkpointing

Create
+
Setup

```
val context = new StreamingContext(...)
val lines = KafkaUtils.createStream(...)
val words = lines.flatMap(...)
...
```



```
def creatingFunc(): StreamingContext = {
  val context = new StreamingContext(...)
  val lines = KafkaUtils.createStream(...)
  val words = lines.flatMap(...)
  ...
  context.checkpoint(hdfsDir)
}
```

Put all setup code into a function that returns a new StreamingContext

Start

```
context.start()
```

Restructuring code for Checkpointing

Create
+
Setup

```
val context = new StreamingContext(...)
val lines = KafkaUtils.createStream(...)
val words = lines.flatMap(...)
...
```



```
def creatingFunc(): StreamingContext = {
  val context = new StreamingContext(...)
  val lines = KafkaUtils.createStream(...)
  val words = lines.flatMap(...)
  ...
  context.checkpoint(hdfsDir)
}
```

Put all setup code into a function that returns a new StreamingContext

Start

```
context.start()
```



```
val context =
  StreamingContext.getOrCreate(
    hdfsDir, creatingFunc)
context.start()
```

Get context setup from HDFS dir OR *create* a new one with the function

Restructuring code for Checkpointing

`StreamingContext.getOrCreate():`

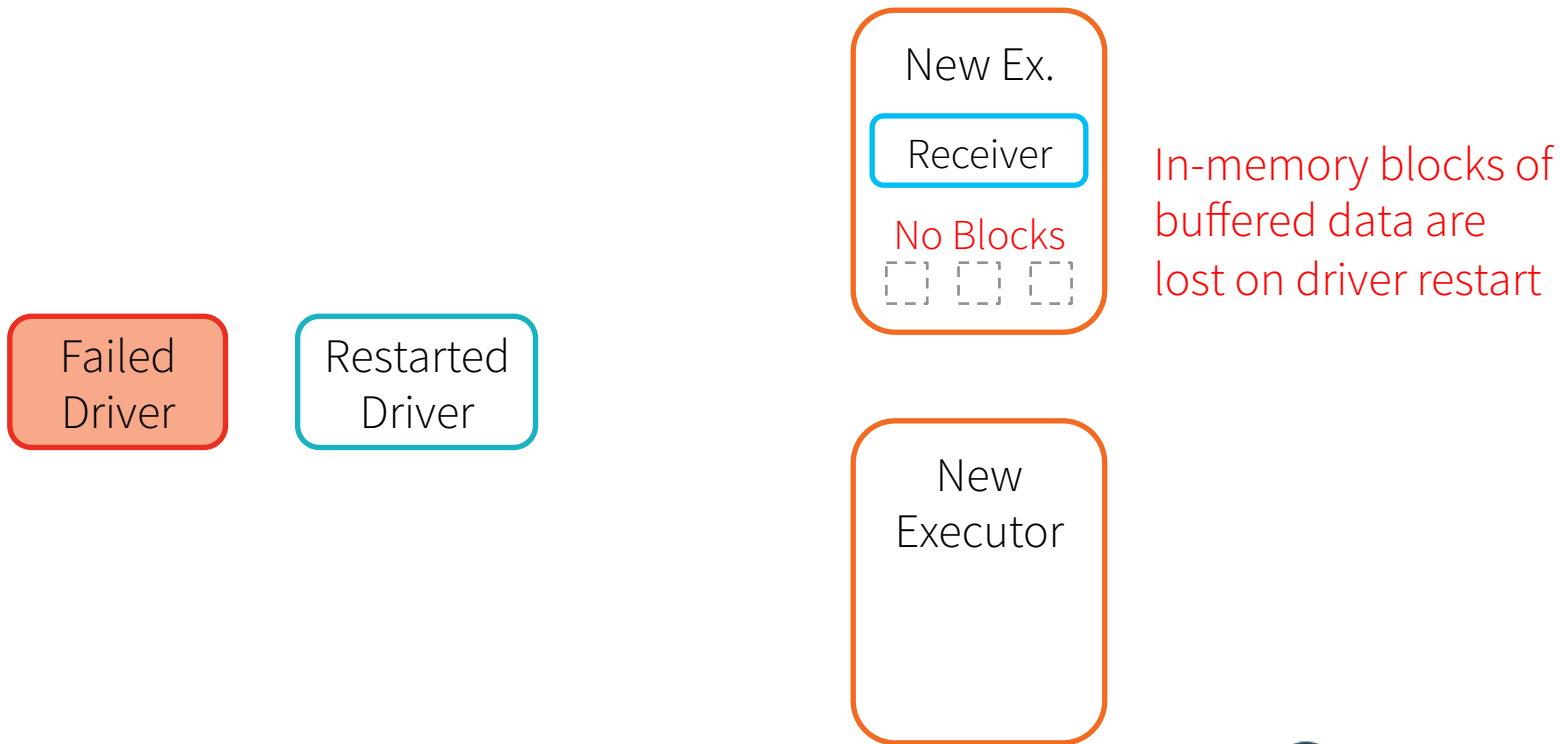
If HDFS directory has checkpoint info
 recover context from info
else
 call `creatingFunc()` to create
 and setup a new context

Restarted process can figure out whether
to recover using checkpoint info or not

```
def creatingFunc(): StreamingContext = {  
  val context = new StreamingContext(...)  
  val lines = KafkaUtils.createStream(...)  
  val words = lines.flatMap(...)  
  ...  
  context.checkpoint(hdfsDir)  
}
```


```
val context =  
  StreamingContext.getOrCreate(  
    hdfsDir, creatingFunc)  
context.start()
```

Received blocks lost on Restart!




Recovering data with Write Ahead Logs

Write Ahead Log (WAL): Synchronously save received data to fault-tolerant storage



Recovering data with Write Ahead Logs

Write Ahead Log (WAL): Synchronously save received data to fault-tolerant storage




Recovering data with Write Ahead Logs

1. Enable checkpointing, logs written in checkpoint directory
3. Enabled WAL in SparkConf configuration
`sparkConf.set("spark.streaming.receiver.writeAheadLog.enable", "true")`
3. Receiver should also be *reliable*
Acknowledge source only after data saved to WAL
Unacked data will be replayed from source by restarted receiver
5. Disable in-memory replication (already replicated by HDFS)
Use `StorageLevel1.MEMORY_AND_DISK_SER` for input DStreams

RDD Checkpointing


- Stateful stream processing can lead to long RDD lineages
- Long lineage = bad for fault-tolerance, too much recomputation
- RDD checkpointing saves RDD data to the fault-tolerant storage to limit lineage and recomputation
- More:
<http://spark.apache.org/docs/latest/streaming-programming-guide.html#checkpointing>

Fault-tolerance Semantics



Zero data loss = every stage processes each event **at least once** despite any failure


Fault-tolerance Semantics



Exactly once, as long as received data is not lost

End-to-end semantics:
At-least once

Fault-tolerance Semantics




Exactly once, as long as received data is not lost

Exactly once, if outputs are idempotent or transactional

End-to-end semantics:
At-least once

Fault-tolerance Semantics




At least once, w/ Checkpointing + WAL + Reliable receivers

Exactly once, as long as received data is not lost

Exactly once, if outputs are idempotent or transactional

End-to-end semantics:
At-least once

Fault-tolerance Semantics



Exactly once receiving with new **Kafka Direct** approach

Treats Kafka like a replicated log, reads it like a file

Does not use receivers

No need to create multiple DStreams and union them


No need to enable Write Ahead Logs

```
val directKafkaStream = KafkaUtils.createDirectStream(...)
```

<https://databricks.com/blog/2015/03/30/improvements-to-kafka-integration-of-spark-streaming.html>

<http://spark.apache.org/docs/latest/streaming-kafka-integration.html>

Fault-tolerance Semantics



Exactly once receiving with new **Kafka Direct** approach

Exactly once, as long as received data is not lost

Exactly once, if outputs are idempotent or transactional

End-to-end semantics:
Exactly once!

Design Tips for Successful Streaming Applications

Areas for consideration

- Enhance resilience with additional components.
- Mini-batch vs. per-message handling.
- Exploit Reactive Streams.

Mini-batch vs. per-message handling

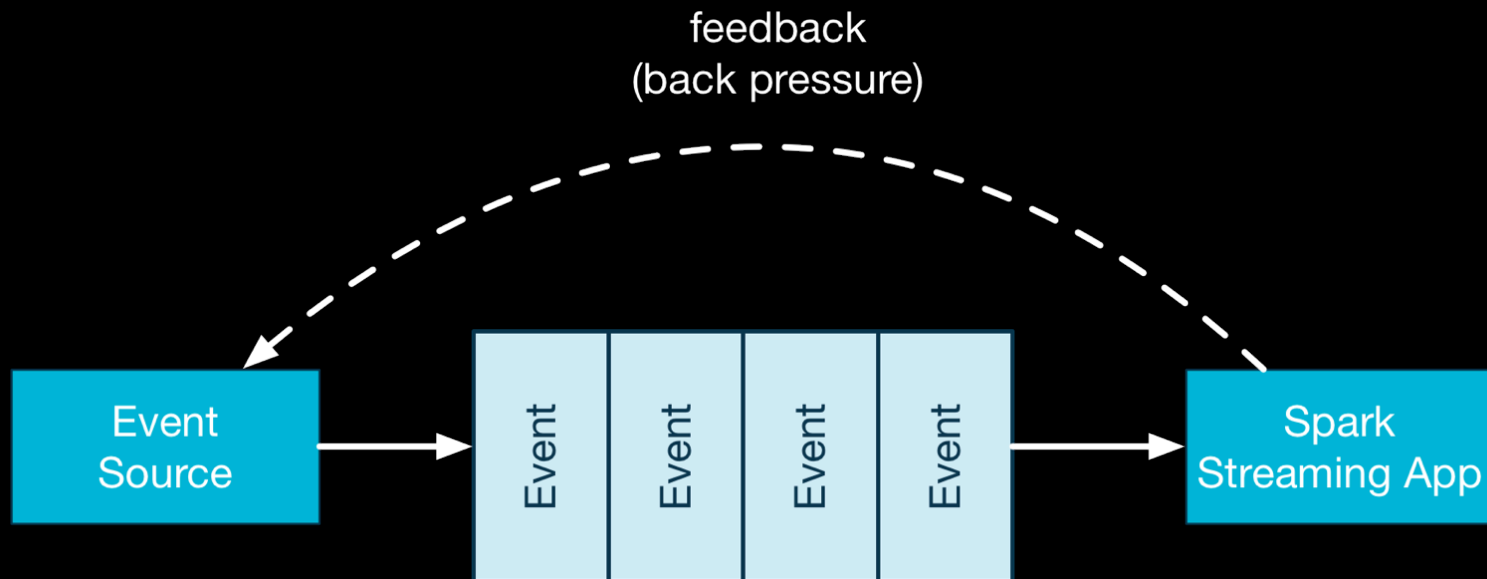
- Use Storm, Akka, Samza, etc. for handling individual messages, especially with sub-second latency requirements.
- Use Spark Streaming's mini-batch model for the Lambda architecture and highly-scalable analytics.

Enhance Resiliency with Additional Components.

- Consider Kafka or Kinesis for resilient buffering in front of Spark Streaming.
 - Buffer for traffic spikes.
 - Re-retrieval of data if an RDD partition is lost and must be reconstructed from the source.
- Going to store the raw data anyway?
 - Do it first, then ingest to Spark from that storage.

Exploit Reactive Streams

- Spark Streaming v1.5 will have support for back pressure to more easily build end-to-end reactive applications

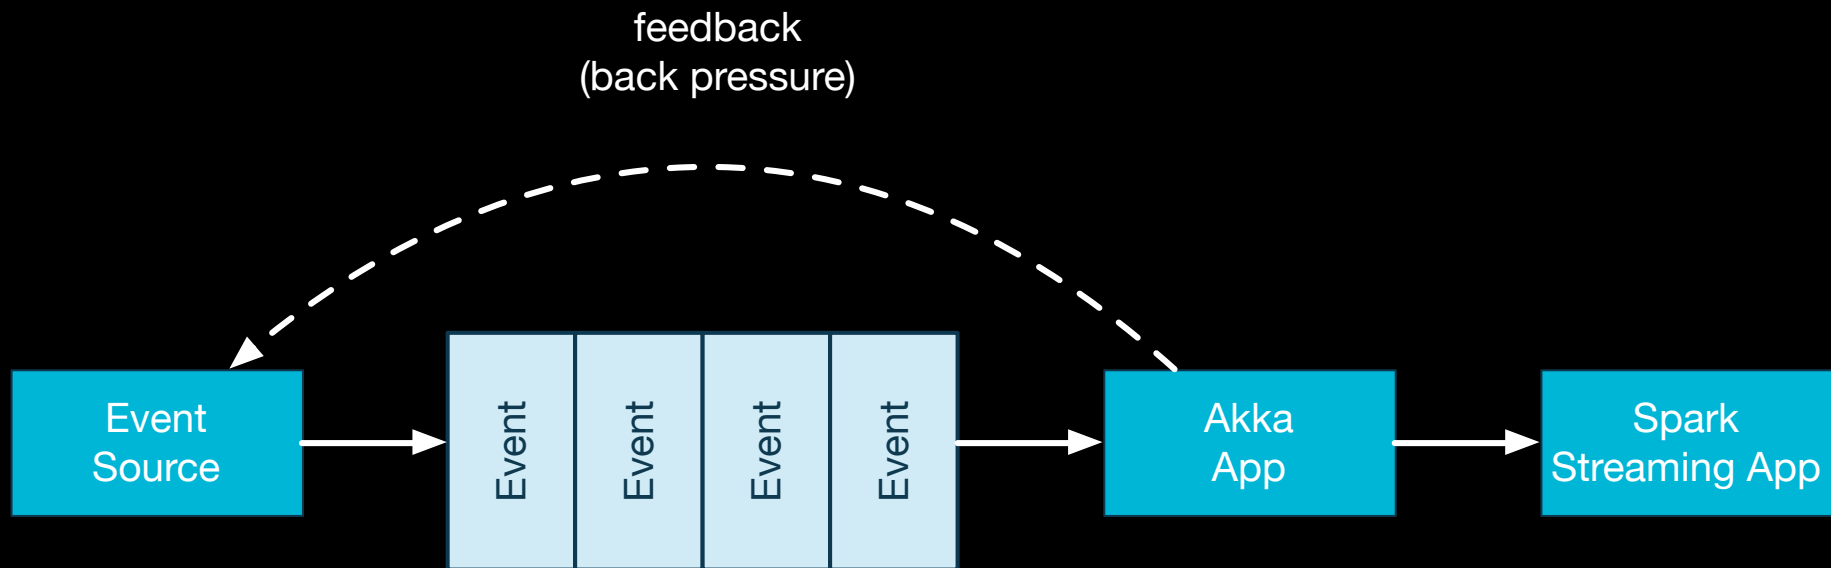


Exploit Reactive Streams

- Spark Streaming v1.5 will have support for back pressure to more easily build end-to-end reactive applications
- Backpressure from consumer to producer:
 - Prevents buffer overflows.
 - Avoids unnecessary throttling.

Exploit Reactive Streams

- Spark Streaming v1.4? Buffer with Akka Streams:



Exploit Reactive Streams

- Spark Streaming v1.4 has a rate limit property:
 - **spark.streaming.receiver.maxRate**
 - Consider setting it for long-running streaming apps with a variable input flow rate.
- Have a graph of Reactive Streams? Consider using an Akka app to buffer the data fed to Spark Streaming over a socket (until 1.5...).

Thank you!

Dean Wampler, Typesafe
Tathagata Das, Databricks



What to do next?

Databricks is available as a hosted platform on AWS with a monthly subscription.

[Start with a free trial](#)

Typesafe now offers certified support for Spark, Mesos & DCOS, read more about it

[READ MORE](#)

REACTIVE PLATFORM

Full Lifecycle Support for Play, Akka, Scala and Spark



Give your project a boost with Reactive Platform:

- **Monitor Message-Driven Apps**
- **Resolve Network Partitions Decisively**
- **Integrate Easily with Legacy Systems**
- **Eliminate Incompatibility & Security Risks**
- **Protect Apps Against Abuse**
- **Expert Support from Dedicated Product Teams**

Enjoy learning? See about the availability of on-site training for Scala, Akka, Play and Spark!

CONTACT US

[Learn more about our offers](#)